

---

# **AREAL: A Large-Scale Asynchronous Reinforcement Learning System for Language Reasoning**

---

**Wei Fu<sup>12</sup>, Jiaxuan Gao<sup>12</sup>, Xujie Shen<sup>2</sup>, Chen Zhu<sup>2</sup>, Zhiyu Mei<sup>12</sup>,  
Chuyi He<sup>2</sup>, Shusheng Xu<sup>12</sup>, Guo Wei<sup>2</sup>, Jun Mei<sup>2</sup>, Jiashu Wang<sup>23</sup>,  
Tongkai Yang<sup>2</sup>, Binhang Yuan<sup>3</sup>, Yi Wu<sup>12</sup>**

<sup>1</sup> IIIS, Tsinghua University, <sup>2</sup> Ant Research, <sup>3</sup> HKUST  
fuwth17@gmail.com, jxwuyi@gmail.com

# Motivation: RL for Advanced Reasoning

- Reinforcement Learning (RL) has become a key paradigm for enhancing the reasoning capabilities of Large Language Models (LLMs).
- RL allows an LLM to generate "thinking tokens" before providing a final answer, improving performance on complex tasks like math, coding, and logic puzzles.
- These models are often called Large Reasoning Models (LRMs).
- Effective RL training requires massive parallelization to generate large batches of sample outputs ("rollouts") for exploration.

# The Problem with Existing Synchronous RL Systems

- Most large-scale RL systems are synchronous, strictly alternating between a "generation" phase and a "training" phase.
- This design ensures training stability, as the model is always trained on the most recent data samples.
- However, this approach suffers from major system-level inefficiencies.
- Primary Bottleneck: The system must wait until the longest output in a generation batch is complete before any model updates can occur.

# GPU Underutilization in Synchronous RL

- The varying lengths of generated responses lead to significant idle time on most GPUs, resulting in poor resource utilization.
- This issue is shown in the timeline below, where GPUs finish at different times but must wait for the slowest one.




Figure 1: Execution timeline of a synchronous (left) and a one-step overlap (right) RL system showing underutilized inference devices.

# Solution: AReaL

- We present AReaL, a fully asynchronous RL system that completely decouples the generation and training processes.
- Key Idea:
  - Rollout (generation) workers continuously generate new data without waiting.
  - Training workers update the model whenever a new batch of data is collected.
  - This design leads to substantially higher GPU utilization and training throughput.

# AReal System Architecture

## Core Components: Workers

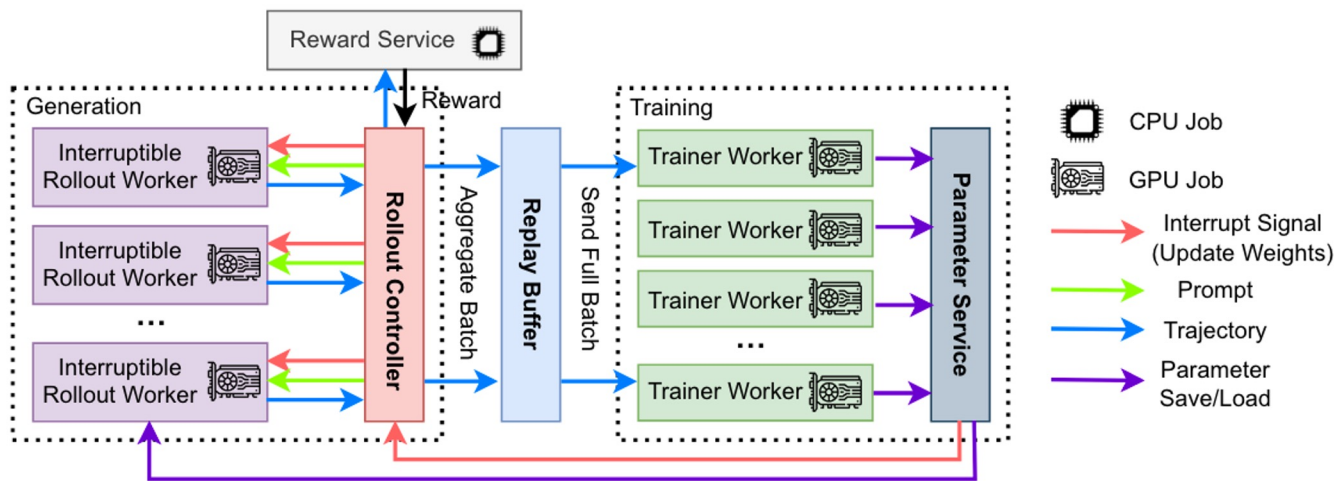


Figure 2: The AREAL architecture featuring asynchronous generation and training components.

# AReal System Architecture

## Core Components: Workers

- **Interruptible Rollout Workers:**
  - Handle requests to generate responses from given prompts.
  - Can be interrupted by `update_weights` requests to load new model parameters mid-generation.
- **Trainer Workers:**
  - Continuously sample data from a Replay Buffer to form a training batch.
  - Perform PPO updates and save the new model parameters.

# AReal System Architecture

## Core Components: Management

- **Rollout Controller:**
  - Acts as the bridge between all components.
  - It sends prompts to rollout workers, forwards completed trajectories to the reward service, and stores the results in the replay buffer.
- **Reward Service:**
  - A separate CPU job that evaluates the correctness of generated responses (e.g., by running unit tests for code).



# Interruptible Generation Workflow

- When new model weights are ready, an interrupt signal is sent.
- Ongoing generations are paused, KV caches from old weights are discarded and recomputed with new weights, and then decoding continues.
- This ensures generation workers are always using the most up-to-date models possible without having to wait for a full batch to complete.




Figure 3: Illustration of generation management in AREAL. Vertical lines show the ready time for the next step training. Blue crosses show the interrupted requests when new parameters arrive.

# Algorithmic Challenges of Asynchronicity

The asynchronous design introduces two main algorithmic challenges:

- **Data Staleness:** Training batches contain data generated from multiple, older policy versions. This can create a distribution gap and degrade learning performance.
- **Inconsistent Policy Versions:** A single trajectory may be generated by segments from different policy versions due to interrupts. This violates the core assumption of the standard PPO algorithm.

# Solution 1: Staleness-Aware Training

- To manage data staleness, we introduce a hyperparameter,  $\eta$ , which defines the maximum permitted staleness for any sample in a training batch.
- When creating a new training batch for policy version  $i$ , we ensure that all data comes from policies no older than  $i - \eta$ .
- The Rollout Controller dynamically controls the rate of generation requests to enforce this constraint.
- This is a simple yet effective way to prevent the model from training on excessively outdated data.

# Solution 2: Decoupled PPO Objective

To handle inconsistent and stale policies, we adopt a **decoupled PPO objective**.

This objective disentangles the **behavior policy** (the policy that generated the data,  $\pi_{\text{behav}}$ ) from the **proximal policy** (the policy used as the baseline for the update,  $\pi_{\text{prox}}$ ).

By using a more recent model as  $\pi_{\text{prox}}$ , we stabilize training by ensuring updates happen within a trust region of a high-quality policy, rather than an old, low-quality one.

This formulation is robust to trajectories generated by multiple policy versions.

$$J(\theta) = \mathbb{E}_{q \sim \mathcal{D}, a_t \sim \pi_{\text{behav}}} \left[ \sum_{t=1}^H \min \left( \underbrace{\frac{\pi_{\theta}}{\pi_{\text{behav}}}}_{\text{Importance Ratio}} \hat{A}_t, \overbrace{\frac{\pi_{\text{prox}}}{\pi_{\text{behav}}} \text{clip} \left( \underbrace{\frac{\pi_{\theta}}{\pi_{\text{prox}}}}_{\text{Trust Region Center}}, 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t}_{\text{Importance Ratio}} \right) \right] \quad (4)$$

$$= \mathbb{E}_{q \sim \mathcal{D}, a_t \sim \pi_{\text{behav}}} \left[ \sum_{t=1}^H \frac{\pi_{\text{prox}}}{\pi_{\text{behav}}} \min \left( u_t^{\text{prox}}(\theta) \hat{A}_t, \text{clip} \left( u_t^{\text{prox}}(\theta), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right) \right], \quad (5)$$

# Experimental Setup

- Tasks: Challenging math (AIME24) and coding (LiveCodeBench) benchmarks.
- Models: Distilled Qwen2 models, with sizes from 1.5B to 32B parameters.
- Hardware: An H800 GPU cluster with up to 64 nodes (512 GPUs).
- Device Allocation: For AReaL, we used a fixed 75-25 split between inference and training devices, respectively, as it yielded the highest throughput in early tests.
- Baselines: We compare against state-of-the-art synchronous systems (DeepScaleR, DeepCoder) and a synchronous variant of AReaL.

# Results: End-to-End Performance

1. AReaL consistently matches or improves final model performance while drastically reducing training time.
2. Across various model sizes, AReaL achieves up to a 2.77x training speedup compared to synchronous systems.

Model	AIME24 ↑	# Nodes	PPO Steps	Training Hours ↓
1.5B basemodel	29.3	-	-	-
w/ VeRL	<b>43.1*</b>	16	250	33.6
w/ Sync.AReaL	42.0	16	250	41.0
w/ AReaL (ours)	42.2	16	250	<b>14.8</b>
7B basemodel	54.3	-	-	-
w/ VeRL	-	24	250	52.1
w/ Sync.AReaL	63.0	24	250	57.7
w/ AReaL (ours)	<b>63.1</b>	24	250	<b>25.4</b>

Model	LiveCodeBench ↑	# Nodes	PPO Steps	Training Hours ↓
14B basemodel	53.4	-	-	-
w/ VeRL	57.9*	32	80	44.4
w/ Sync.AReaL	56.7	32	80	48.8
w/ AReaL (ours)	<b>58.1</b>	32	80	<b>21.9</b>
32B basemodel	57.4	-	-	-
w/ VeRL	-	48	60	46.4
w/ Sync.AReaL	<b>61.2</b>	48	60	51.1
w/ AReaL (ours)	61.0	48	60	<b>31.1</b>

# Results: Scalability

1. We compared the strong-scaling of AReal against verl, a state-of-the-art synchronous system.
2. AReal demonstrates nearly linear scaling as the number of GPUs increases.
3. The synchronous system fails to scale effectively, especially with longer context lengths.


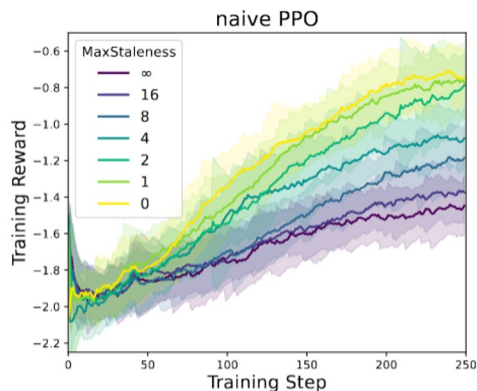



Figure 4: The strong scaling trend. Dotted lines indicate ideal linear scaling. verl consistently encounters OOM with 32k context length and the 32B model so the data points are missing.

# Algorithm Ablation: Staleness and Decoupled PPO

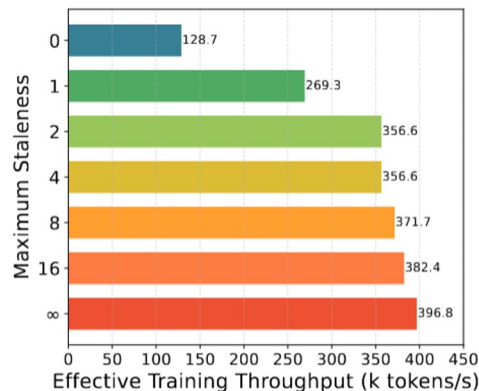
1. Naive PPO (left): Performance degrades significantly as data staleness ( $\eta$ ) increases.
2. Decoupled PPO (center): The decoupled objective substantially improves training stability and performance, even with stale data.
3. Throughput (right): Allowing for moderate staleness dramatically increases effective training throughput.



(a) Learning curves with naive PPO.



(b) Learning curves with eq. (5).



(c) Effective training throughput.



# Algorithm Ablation: Performance vs. Staleness

1. With the decoupled objective, a moderate maximum staleness ( $\eta=4$  or  $\eta=8$ ) achieves performance comparable to the synchronous "oracle" ( $\eta=0$ ).
2. However, unbounded staleness still leads to inferior performance.
3. This validates our approach of combining controlled staleness with the decoupled PPO objective.

Max.Stale.	AIME24		AIME25		AMC23		MATH 500	
	W/o	With	W/o	With	W/o	With	W/o	With
0 (Oracle)	42.0		32.9		84.4		89.2	
1	<u>41.8</u>	<u>42.1</u>	30.7	<u>31.9</u>	<u>83.3</u>	<u>85.2</u>	<u>89.9</u>	<u>89.8</u>
2	40.0	<u>41.8</u>	<u>32.1</u>	<u>32.5</u>	82.3	<u>84.3</u>	<u>89.6</u>	<u>89.6</u>
4	23.3	<u>42.2</u>	23.1	<u>32.0</u>	58.5	<u>85.1</u>	66.9	<u>89.5</u>
8	35.7	<u>41.0</u>	27.8	<u>31.1</u>	81.2	82.9	87.8	<u>89.2</u>
16	35.8	38.7	26.2	<u>32.5</u>	78.4	83.2	87.4	<u>89.1</u>
$\infty$	34.0	36.9	26.9	<u>29.9</u>	79.4	81.0	87.1	88.1

# System Ablation: Interruptible Generation

1. We compared the throughput of our system with and without the interruptible generation feature.
2. Interruptible generation leads to a 12% throughput increase for the 1.5B model and a 17% increase for the 7B model.
3. This confirms that dynamically updating weights without waiting for slow responses to finish is a key architectural benefit.




Figure 6: Ablation study of interruptible generation.

# System Ablation: Dynamic Batching

1. We evaluated our dynamic micro-batch allocation algorithm against a standard strategy.
2. Our algorithm intelligently balances tokens across micro-batches to maximize GPU memory utilization and minimize padding.
3. Dynamic batching yields an average of 30% throughput improvement during PPO training across all tested model sizes.




Figure 7: Ablation study of dynamic micro-batch allocation.

# Conclusion

1. We introduced AReaL, a fully asynchronous system for large-scale RL training that is efficient, scalable, and stable.
2. By completely decoupling generation and training, AReaL achieves superior hardware utilization and up to a 2.77x training speedup.
3. Key Innovations:
  - a. An expressive, asynchronous architecture with interruptible workers.
  - b. Algorithmic enhancements—staleness-aware training and a decoupled PPO objective—that stabilize training with stale data.
4. This work provides a robust foundation for reliably scaling RL, enabling future advances in machine intelligence.